

Why compare

Toy estimators need a harness: same JSON in, overflow and runtime out

The idea

- Fix the instance
- Run RUDY and probabilistic
- Record total overflow, max overflow, wall time
- Optionally compare to a reference map
- Write one discrepancy hypothesis

Harness shape

- A small Python driver loading `tiny_cong.json` is enough
- Print a markdown table
- Missing external tools are OK, document blockers and still validate I/O

Pitfalls

- Changing Cap between rows of the table
- Comparing incompatible demand units
- Optimizing runtime before correctness

Your turn

- Fill one comparison table for spread versus congested seeds
- Then open the wrap